

Lecture 8: Deep Reinforcement Learning

CS60077 : REINFORCEMENT LEARNING


Autumn 2023

Outline

- 1 Introduction to Deep RL
- 2 Deep Q-Networks
- 3 Deep RL in Atari Games
- 4 DQN Extensions

Deep Neural Networks as Function Approximators

- Use Deep Neural Networks (DNN) to represent
 - State Value (V) Function
 - State-Action Value (Q) Function
 - Parameterized Policy
 - Model
- Optimize Loss Function by Stochastic Gradient Descent



Q-Learning with Value Function Approximation

- Q-learning under Tabular Representation:
 - Converges to the optimal $Q^*(s, a)$
- Q-learning with Value Function Approximation (VFA):
 - Minimize MSE loss by stochastic gradient descent using a target Q estimate instead of true Q value
 - Q-learning with VFA can *diverge*
- Two of the issues causing problems:
 - **Correlations** between Samples
 - Non-Stationary Targets
- Deep Q-learning (DQN) addresses these challenges by
 - Experience Replay (removing correlations)
 - Fixed Q-Targets (improving stability)

DQNs with Experience Replay

- To help remove correlations, store dataset \mathcal{D} (called a **replay buffer**) from prior experience

s_1, a_1, r_1, s_2	$\rightarrow s, a, r, s'$
s_2, a_2, r_2, s_3	
s_3, a_3, r_3, s_4	
\dots	
$s_t, a_t, r_{t+1}, s_{t+1}$	

- To perform experience replay, repeat the following:
 - $(s, a, r, s') \sim \mathcal{D}$: sample an experience tuple from the dataset
 - Compute the target value for the sampled

$$s : r + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s', a'; \mathbf{w})$$

- Use stochastic gradient descent to update the network weights

$$\Delta \mathbf{w} = \alpha \left(r + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s', a'; \mathbf{w}) - \hat{Q}(s, a; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

- Uses target as a scalar, but function weights will get updated on the next round, changing the target value

DQNs with Fixed Q-Targets

- To help improve stability, fix the **target weights** used in the target calculation for multiple updates
- Target network uses a different set of weights than the weights being updated
- Let parameters \mathbf{w}^- be the set of weights used in the target, and \mathbf{w} be the weights that are being updated
- Slight change to computation of target value:
 - $(s, a, r, s') \sim \mathcal{D}$: sample an experience tuple from the dataset
 - Compute the target value for the sampled

$$s : r + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s', a'; \mathbf{w}^-)$$

- Use stochastic gradient descent to update the network weights


$$\Delta \mathbf{w} = \alpha \left(r + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s', a'; \mathbf{w}^-) - \hat{Q}(s, a; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a; \mathbf{w})$$

DQN Strategy

DQN uses **experience replay** and **fixed Q-targets**

- Take action a_t according to ϵ -greedy policy
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory \mathcal{D}
- Sample random mini-batch of transitions (s, a, r, s') from \mathcal{D}
- Compute Q-learning targets w.r.t. old, fixed parameters \mathbf{w}^-
- Optimise MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(w_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}_i} \left[\left(r + \gamma \max_{a \in \mathcal{A}} \hat{Q}(s', a'; w_i^-) - Q(s, a; w_i) \right)^2 \right]$$
- Using variant of stochastic gradient descent



DQN Algorithm

Algorithm 1: Deep Q-Learning with Experience Replay and Fixed Q-Targets


```

1 Initialize replay memory  $\mathcal{D}$  to capacity  $N$ ;
2 Initialize action-value function  $q$  with random weights;
3 Initialize target action-value function  $\hat{Q}$  with weights  $\mathbf{w}^- = \mathbf{w}$ ;
4 for  $episode = 1, M$  do
5     Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$ ;
6     for  $t = 1, T$  do
7         Select  $a_t = \begin{cases} \text{random action,} & \text{with probability } \epsilon \\ \arg \max_{a \in \mathcal{A}} Q(\phi(s_t), a; \mathbf{w}), & \text{otherwise} \end{cases}$ ;
8         Execute action  $a_t$  and observe reward  $r_t$  and image  $x_{t+1}$ ;
9         Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ ;
10        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$ ;
11        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$ ;
12        Set  $y_j = \begin{cases} r_j, & \text{if episode terminates at step } (j + 1) \\ r_j + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(\phi_{j+1}, a', \mathbf{w}^-), & \text{otherwise} \end{cases}$ ;
13        Perform gradient descent over  $(y_j - Q(\phi_j, a_j; \mathbf{w}))^2$  w.r.t. parameter  $\mathbf{w}$ ;
14        After every  $C$  steps, reset  $\mathbf{w}^- \leftarrow \mathbf{w}$ ;
15    end
16 end

```

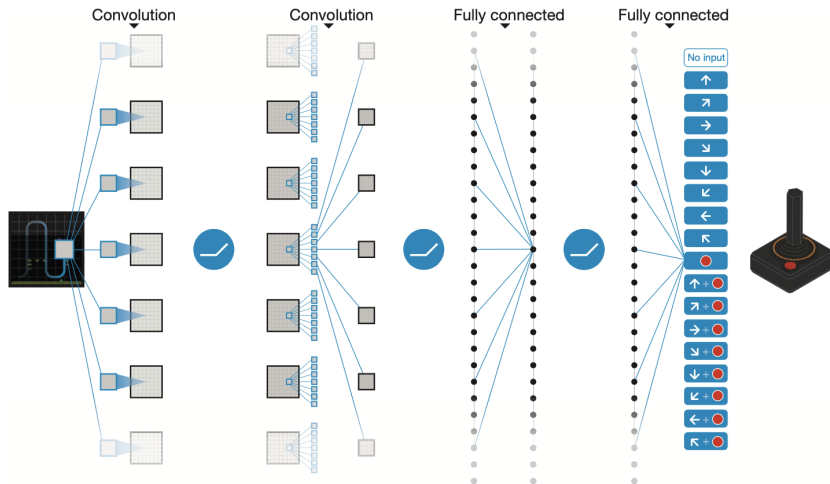

DQN in Atari Games

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is stack of raw pixels from last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step




(Network architecture and hyperparameters fixed across all games)

Human-level Control through DQN Models



DQN Results in Atari Games



Effects of Replay and Separating Target Q-network

Atari	DQN with Replay		DQN without Replay		Linear
	with Fixed Target Q	without Target Q	with Fixed Target Q	without Target Q	
Breakout	316.8	240.7	10.2	3.2	3.0
Enduro	1006.3	831.4	141.9	29.1	62.0
River Raid	7446.6	4102.8	2867.7	1453.0	2346.9
Seaquest	2894.4	822.6	1003.0	275.8	656.9
Space Invaders	1088.9	826.3	373.2	302.0	301.3

Ref: Mnih et. al.; "Human-Level Control through Deep Reinforcement Learning"; Nature, 518:529-533, 2015.

Improvements over DQN Framework

- Double DQN (DDQN)
- Prioritized Replay DDQN
- Dueling DQN
- Distributional DQN
- Noisy DQN

Double Deep Q-Networks (DDQN)

- Q-learning update for the (weight) parameters:

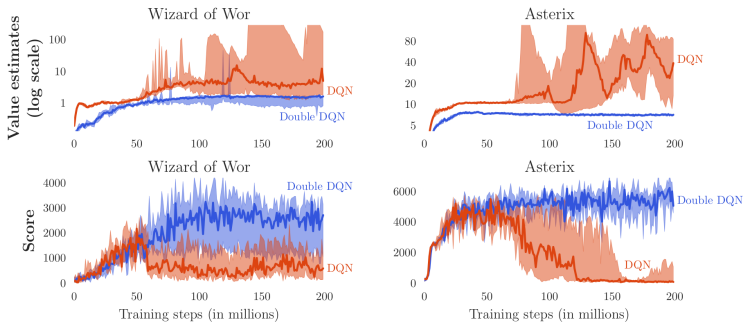
$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha(y_t - Q(s_t, a_t; \mathbf{w}_t)) \nabla_{\mathbf{w}_t} Q(s_t, a_t; \mathbf{w}_t)$, where:

- (Standard) Q-learning target: $y_t = r_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a; \mathbf{w}_t)$
 - DQN target: $y_t = r_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a; \mathbf{w}_t^-)$
- Issue: There is an upward bias in $\max_{a \in \mathcal{A}} Q(s, a; \mathbf{w})$ estimates
- Double DQN maintains two sets of weight \mathbf{w} and \mathbf{w}^- , so reduce bias by using:
 - \mathbf{w} for selecting the best action
 - \mathbf{w}^- for evaluating the best action
- Double DQN loss:

$$\mathcal{L}_i(w_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}_i} \left[\underbrace{\left(r + \gamma Q(s', \arg \max_{a' \in \mathcal{A}} Q(s', a'; w_i), w_i^-) - Q(s, a; w_i) \right)^2}_{\text{Double DQN Target}} \right]$$

Performance of Double DQN

- Value estimates and scores by DQN and Double DQN on 6 Atari games




- Summary of normalized performance on 49 Atari games with human starts

	upto 5 min of play		upto 30 min of play		
	DQN	DDQN	DQN	DDQN	DDQN (tuned)
Median	93.5%	114.7%	47.5%	88.4%	116.7%
Mean	241.1%	330.3%	122.0%	273.1%	475.2%

Double DQN in Atari Games

Normalized scores on 57 Atari games, tested for 100 episodes per game with human starts



Prioritized Experience Replay (with DDQN)

- Replaying all transitions with equal probability is highly suboptimal
- Replay transitions in proportion to absolute Bellman error:


$$\left| r + \gamma \max_{a' \in \mathcal{A}} Q(s', a'; \mathbf{w}^-) - Q(s, a; \mathbf{w}) \right|$$

- Leads to much faster learning
- Prioritized replay to DQN adds substantial improvements in (normalized) score on 41 out of 49 Atari games

	Double DQN		Double DQN (tuned)		
	Baseline	Rank-based	Baseline	Rank-based	Proportional
Median	48%	106%	111%	113%	128%
Mean	122%	355%	418%	454%	551%
> Baseline	—	41	—	38	42
> Human	15	25	30	33	33
# Games	49	49	57	57	57

Prioritized Replay DDQN in Atari Games

- Substantial improvements in most games found through difference in normalized score (gap between random and human is 100%) on 57 games with human starts
- Comparing Double DQN with and without prioritized replay (**rank-based** as well as **proportional** variant)




Dueling Deep Q-Networks

Value-Advantage Decomposition of Q:

$$Q(s, a; \mathbf{w}, \alpha, \beta) = V(s; \mathbf{w}, \beta) + A(s, a; \mathbf{w}, \alpha), \text{ where}$$

- \mathbf{w} denotes convolutional layer parameters
- α and β are the FC-layer parameters of two streams



- *Identifiability Issue*: Given Q , cannot recover V and A uniquely
- To address this issue of identifiability, force advantage function estimator to have zero advantage at chosen action

$$Q(s, a; \mathbf{w}, \alpha, \beta) = V(s; \mathbf{w}, \beta) + \left(A(s, a; \mathbf{w}, \alpha) - \max_{a' \in |\mathcal{A}|} A(s, a'; \mathbf{w}, \alpha) \right)$$


$$Q(s, a; \mathbf{w}, \alpha, \beta) = V(s; \mathbf{w}, \beta) + \left(A(s, a; \mathbf{w}, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a' \in |\mathcal{A}|} A(s, a'; \mathbf{w}, \alpha) \right)$$

Dueling DQN in Atari Games

- Mean and median scores across all 57 Atari games, measured in percentages of human performance


DQN Variant	30 no-ops		Human Starts	
	Mean	Median	Mean	Median
Prior. Duel Clip	591.9%	172.1%	567.0%	115.3%
Prior. Single	434.6%	123.7%	386.7%	112.9%
Duel Clip	373.1%	151.5%	343.8%	117.1%
Single Clip	341.2%	132.6%	302.8%	114.1%
Single	307.3%	117.8%	332.9%	110.9%
Nature DQN	227.9%	79.1%	219.6%	68.5%

- Improvements of dueling architecture over Prioritized DDQN baseline




Ref: Wang et. al.; "Dueling Network Architectures for Deep Reinforcement Learning"; PMLR, 48:1995-2003, 2016.


Dueling DQN Performance in Atari Games



(a) Improvement in percentage of NoisyNet-DQN over DQN



(b) Improvement in percentage of NoisyNet-Dueling over Dueling



(c) Improvement in percentage of NoisyNet-A3C over A3C


Distributional Bellman Equations and Updates (no exam)

- Instead of learning expected return $Q(s, a)$, learn to estimate distribution of returns whose expectation is $Q(s, a)$
- The value function, $Q^\pi(s, a) = \mathbb{E}[Z^\pi(s, a)]$, where $Z^\pi(s, a) = \sum_{t \geq 0} \gamma^t R(s_t, a_t)$ satisfies Bellman equation $Q^\pi(s, a) = R(s, a) + \gamma Q^\pi(s', a')$. Z is the support of the return probability distribution (a categorical distribution)

- Distributional Bellman Equation:

$$Z^\pi(s, a) = R(s, a) + \gamma Z^\pi(s', a'), \text{ where } s' \sim p(\cdot | s, a) \text{ and } a' \sim \pi(\cdot | s')$$

- Distributional Bellman Operator: $T^\pi Z(s, a) = R(s, a) + \gamma Z(s', a')$



- Distributional Bellman Optimality Operator:

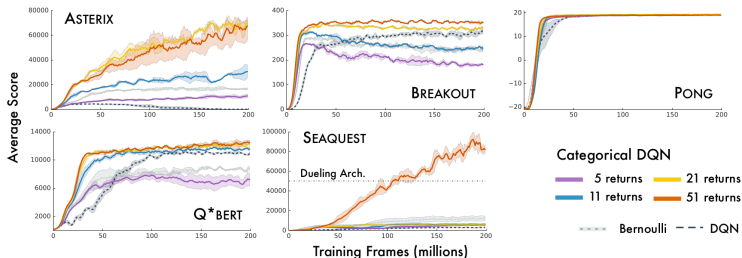
$$TZ(s, a) = R(s, a) + \gamma Z(s', \pi_Z(s')), \text{ where } \pi_Z(s') = \arg \max_{a' \in \mathcal{A}} \mathbb{E}[Z(s', a')]$$

Categorical/Distributional DQN Performance (no exam)

- Mean and median scores across all 57 Atari games, measured in percentages of human performance

	DQN	Double DQN	Dueling DQN	Prior. DQN	Prior.Duel. DQN	Distributional DQN (C521)	UNREAL
Mean	228%	307%	373%	434%	592%	701%	880%
Median	79%	118%	151%	124%	172%	178%	250%
> Human	24	33	37	39	39	40	—
> DQN	0	43	50	48	44	50	—

- Categorical DQN: Varying number of atoms in the discrete distribution. Scores are moving averages over 5 million frames.



Noisy Nets for Exploration

(no exam)

- Instead of ϵ -greedy policy, Noisy Nets inject noise in the parametric space for exploration.
- Add noise to network parameters for better exploration
 - Use factorised Gaussian noise for DQN and Duelling
 - Use independent Gaussian noise for actor-critic method (distributed A3C)
- Noisy Nets are neural networks whose weights and biases are perturbed by parametric function of noise
 - *Standard Linear Layer*: $y = wx + b$, where $x \in \mathbb{R}^p$ is the layer input, $w \in \mathbb{R}^{q \times p}$ the weight matrix, and $b \in \mathbb{R}^q$ the bias
 - *Noisy Linear Layer*: $y = (\mu^w + \sigma^w \odot \epsilon^w)x + (\mu^b + \sigma^b \odot \epsilon^b)$, where $\mu^w \in \mathbb{R}^{q \times p}$, $\mu^b \in \mathbb{R}^q$, $\sigma^w \in \mathbb{R}^{q \times p}$ and $\sigma^b \in \mathbb{R}^q$ are learnable, whereas $\epsilon^w \in \mathbb{R}^{q \times p}$ and $\epsilon^b \in \mathbb{R}^q$ are random noise variables




Noisy Nets in Atari Games

(no exam)

- Mean and median scores across all 57 Atari games, measured in percentages of human performance

Method	Baseline		Noisy Net		Improvement (On Median)
	Mean	Median	Mean	Median	
DQN	319	83	379	123	48%
Dueling DQN	524	132	633	172	30%
A3C	293	80	347	94	18%

- Comparison of learning curves of NoisyNet agent versus baseline according to median human normalised score



Rainbow: The Integrated DQN

(no exam)

- *Distributional RL* estimate for return distribution than expected returns
- Replace 1-step distributional loss with *Multi-step Variant*
- This corresponds to a target distribution,
 $d_t^{(n)} = (R_t^{(n)} + \gamma_t^{(n)} \mathbf{z}, p_{\bar{\theta}}(s_{t+n}, \mathbf{a}_{t+n}^*))$, with the KL-loss as $D_{KL}(\Phi_{\mathbf{z}} d_t^{(n)} || d_t)$,
 where $\Phi_{\mathbf{z}}$ is a projection onto \mathbf{z}
- Combine distributional loss with multi-step *Double Q-Learning*
- Incorporate concept of *Prioritized Replay* by sampling transitions which are prioritized by KL loss
- Network architecture is a *Dueling Architecture* adapted for use with return distributions
- For each atom \mathbf{z}^i , corresponding probability mass $p_{\theta}^i(s, a)$ is calculated as below by using a dueling architecture:

$$p_{\theta}^i(s, a) \propto \frac{\exp(v_n^i(f_{\xi}(s)) + a_{\psi}^i(f_{\xi}(s), a) - \frac{1}{|\mathcal{A}|} \sum_{a'} a_{\psi}^i(f_{\xi}(s), a'))}{\sum_j \exp(v_n^j(f_{\xi}(s)) + a_{\psi}^j(f_{\xi}(s), a) - \frac{1}{|\mathcal{A}|} \sum_{a'} a_{\psi}^j(f_{\xi}(s), a'))}$$

- Finally incorporate noisy streams by *replacing all linear layers with their noisy equivalent*

Rainbow Performance Results


(no exam)

- Median normalized scores of the best agent snapshots for Rainbow and baselines

Agent	no-ops	human starts
DQN	79%	68%
DDQN	117%	110%
Prio. DQN	140%	128%
Duel. DQN	151%	117%
A3C	—	116%
Noisy DQN	118%	102%
Dist. DQN	164%	125%
Rainbow	223%	153%


- Demerit:** *Computationally intensive* – takes 10 days for 200M frames (full run) on 1 GPU (so 57 games would take 570 days on a single GPU)

- Median human-normalized performance across 57 Atari games, as a function of time



Rainbow Performance over Atari Games

(no exam)



Ref: Hessel et. al.; "Rainbow: Combining Improvements in Deep Reinforcement Learning"; AAAI, 32(1):3215-3222, 2018.

Thank You!

Questions?

The only stupid question is the one you were afraid to ask but never did!
– Rich Sutton